



Accelerating Coverage Closure with Reinforcement Learning: A Case Study on FSM Verification

Tijana Misic, Verification Consulting



Agenda

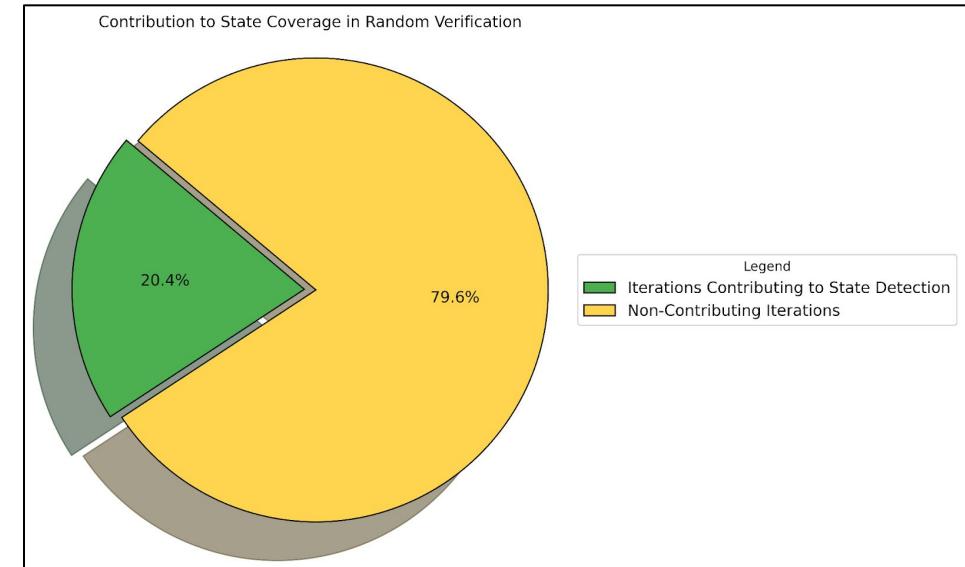
- Motivation
- Proposed Approach
- Implementation Results
- ROI & Discussion
- Conclusion

Motivation

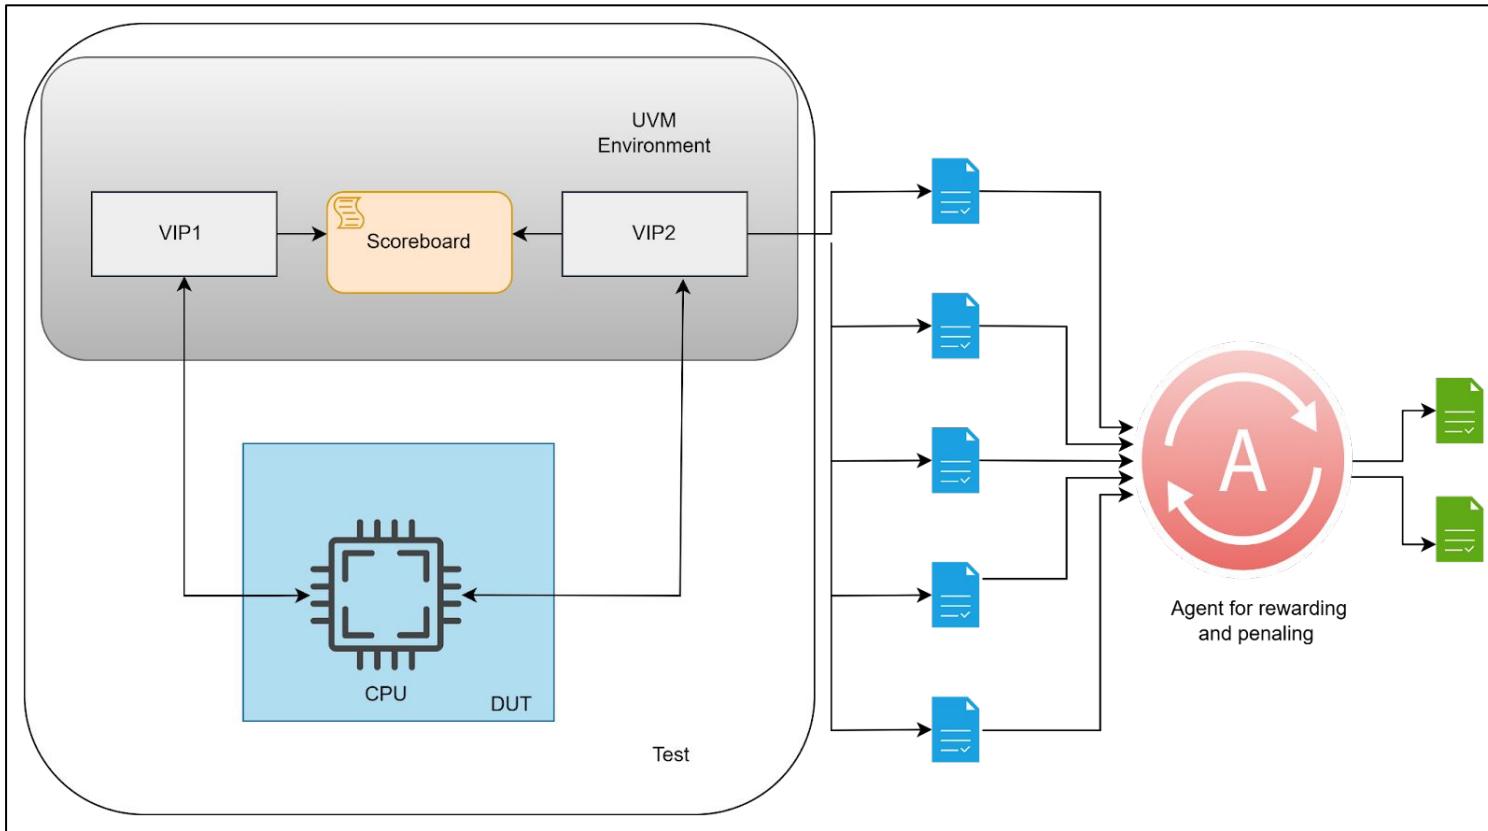
- Functional coverage closure is a bottleneck
- Random regressions inefficient
- Need intelligent guidance
- Reinforcement Learning provides adaptive optimization

Proposed Approach

- RL Agent guides stimulus selection
- Reward = new coverage
- Penalty = redundant test
- Iterative loop until closure



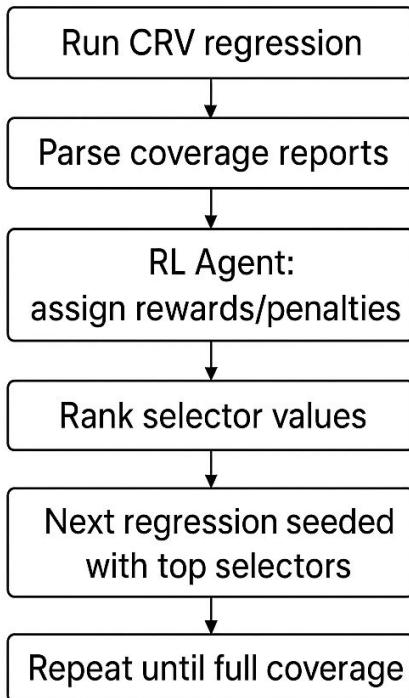
UVM Testbench and RL Agent Integration



Implementation

- Agent in Python
- Minimal integration
- One random test with 3 sets of constraints
 - sel[0] - External events (Wake-up, sleep)
 - sel[1] - OverVoltage, UnderVoltage
 - sel[2] - Others

RL-Augmented Verification Flow



Pseudocode

```
for iteration in CRV_regression:  
    run_test(iteration)  
    covered_states = pars_coverage_report()  
    if new_state_covered(covered_states):  
        reward(iteration)  
    else:  
        penalize(iteration)  
  
    # Rank selector values based on rewards  
    top_selectors = rank_selectors_by_reward()  
  
    # Seed next regression with top selectors  
    for selector in top_selectors:  
        run_test_with_selector(selector)
```

Result of RL Agent processing

- The Part of Test Ranking Report

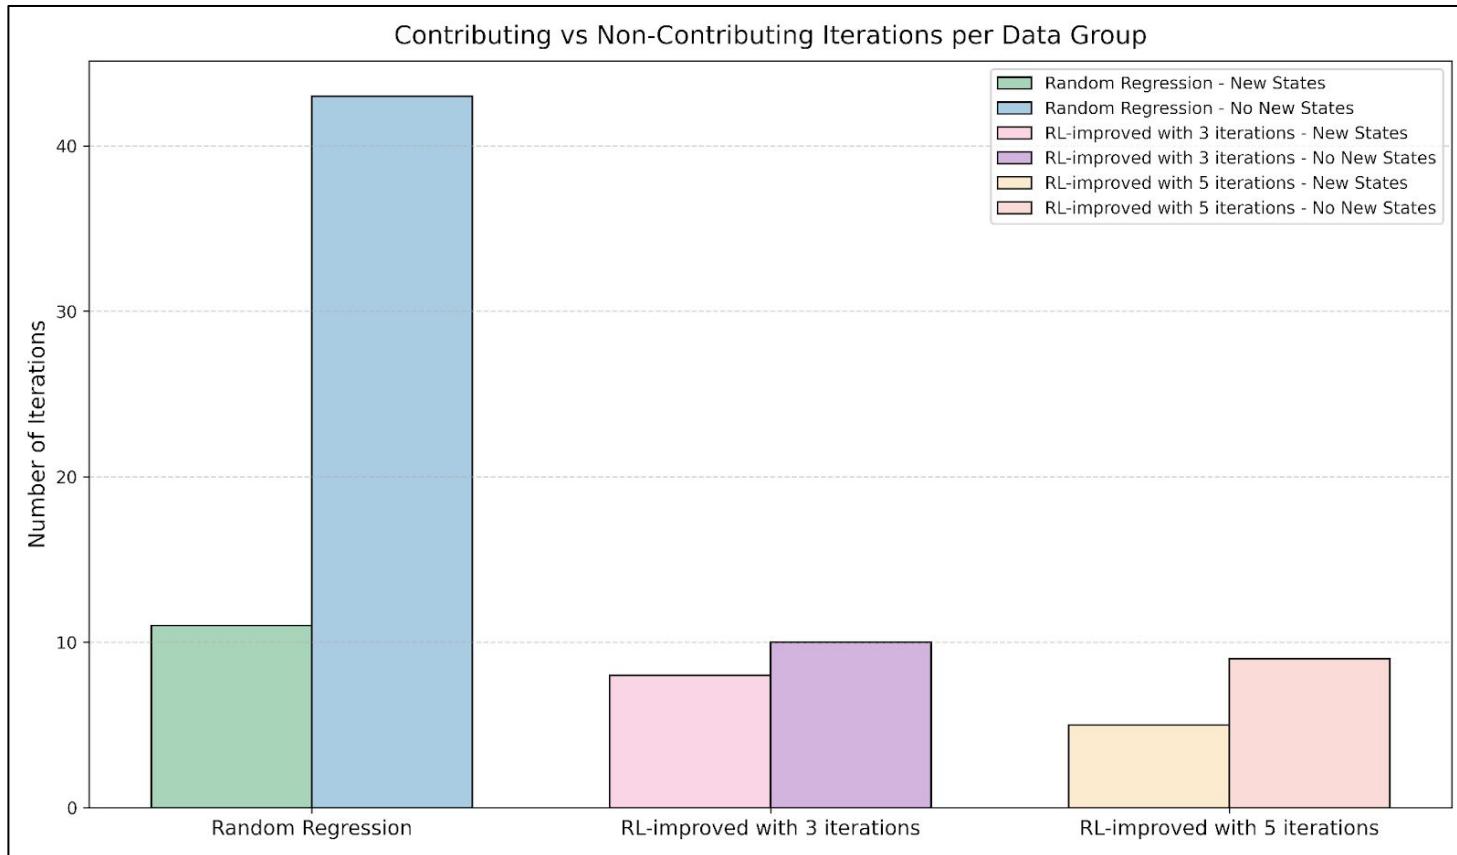
[NEW]	[TEST_1]	State: 0000,1010	Constraints: [0, 0, 0, 0] => Reward: 1.0
[NO_NEW_COV]	[TEST_2]	State: /	Constraints: [0, 1, 1, 1] => Penalty: -0.5
[NEW]	[TEST_3]	State: 0100,1000,1011,1100,1111	Constraints: [0, 0, 0, 0] => Reward: 1.0
[NEW]	[TEST_4]	State: 1001	Constraints: [0, 1, 1, 1] => Reward: 1.0
[NO_NEW_COV]	[TEST_5]	State: /	Constraints: [0, 1, 1, 1] => Penalty: -0.5
[NO_NEW_COV]	[TEST_6]	State: /	Constraints: [0, 0, 1, 0] => Penalty: -0.5
[NEW]	[TEST_7]	State: 0001	Constraints: [0, 1, 0, 1] => Reward: 1.0
[NEW]	[TEST_8]	State: 1101	Constraints: [0, 1, 0, 1] => Reward: 1.0
[NO_NEW_COV]	[TEST_9]	State: /	Constraints: [0, 1, 0, 1] => Penalty: -0.5

- An optimized subset of constraints that maximizes coverage efficiency:
`b000, `b111, `b101, `b011, `b001

Results

- FSM design experiment
- RL reduces regressions
- RL achieves ~3x faster coverage closure on the FSM benchmark
- Single random test run is called iteration = effort unit

Efficiency Gains with RL over CRV



ROI and Discussion

- Measure = iterations (number of test runs) to closure
- Faster closure means fewer regression cycles and reduced compute cost.
- Low overhead
- Reduced regression cycles translate directly to shorter verification turnaround and compute savings.
- ROI = fewer iterations + lower compute cost

Future Work

- The reward model can be extended to include:
 - Test runtime - to prefer faster tests
 - Critical features - to prioritize tests which are covering critical features
 - Diversity- to encourage test variety
- Evaluation of the proposed flow to more complex verification environments and designs

Conclusion

- RL accelerates coverage closure
- FSM case study
- Generalizable methodology
- Future: larger designs
- Key takeaway: fewer iterations = faster closure

Questions

- Thank you for your attention
- All questions and suggestions are welcome

